

Why internal-state probing: $\sim 300\times$ cheaper at inference

Detection latency per response (s, log scale)

