

Inference cost: best configurations vs. baselines (six-billion model, single GPU)

